



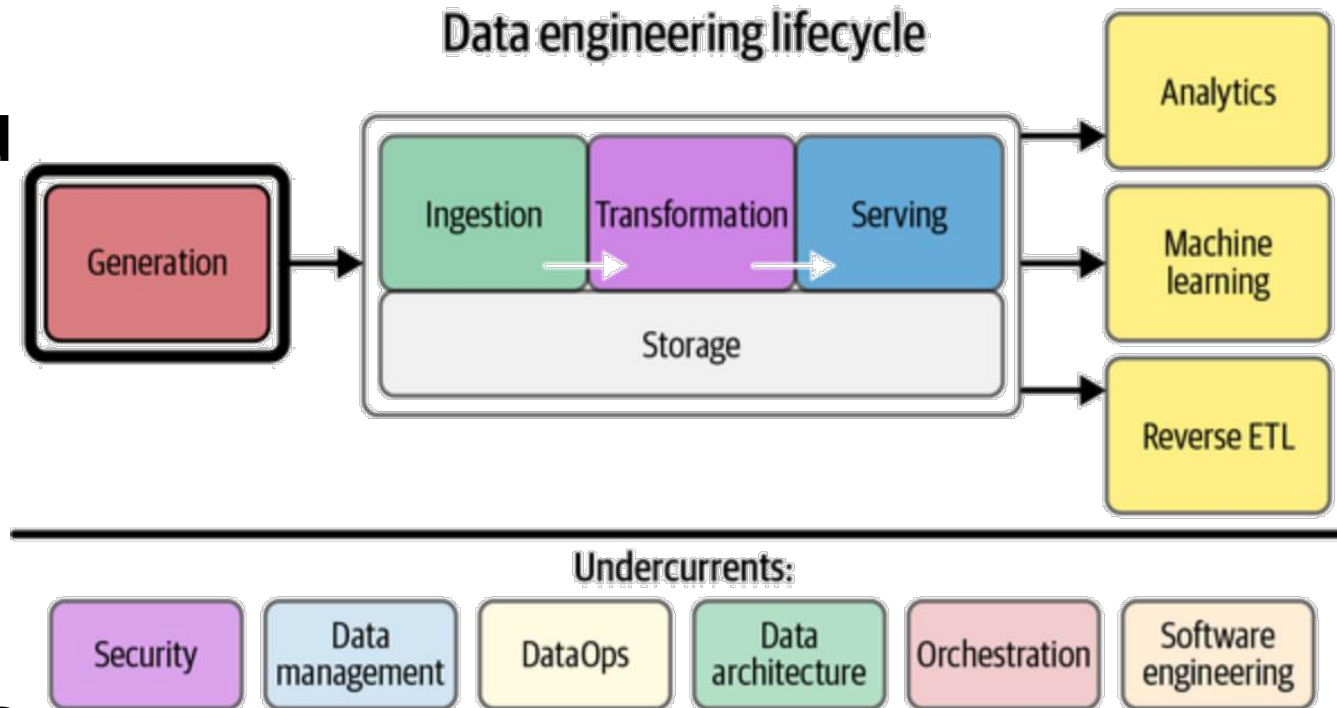
Fundamentals of Data Engineering

Trainer: Pradnyaa S Dindorkar



Data engineering

- Data engineering is the development, implementation, and maintenance of systems and processes that take in raw data and produce high-quality, consistent information that supports downstream use cases, such as analysis and machine learning.
- Data engineer manages data engineering lifecycle, beginning with getting data from source systems & ending with serving data for use cases, such as analysis or machine learning.



https://youtu.be/hZu_87l62J4



Traditional ETL vs Hadoop ELT

- **ETL stands for Extract, Transform and Load.**
- **The ETL process typically extracts data from the source/transactional systems, transforms it to fit the model of data-warehouse and finally loads it to the data warehouse.**
- **The transformation process involves cleansing, enriching and applying transformations to create desired output.**
- **Data is usually dumped to a staging area after extraction.**
- **ELT stands for Extract, Load and Transform.**
- **As opposed to loading just the transformed data in the target systems, the ELT process loads the entire data into the data lake. This results in faster load times.**
- **The load process can also perform some basic validations and data cleansing rules.**
- **The data is then transformed for analytical reporting as per demand.**



Data storage


- **Data storage is related to multiple stages in data engineering life cycle i.e. ingestion, transformation and serving.**
- **Storage needs to be selected based on read/write requirement, speed, durability, consistency, availability, scalability, fault tolerance, ... factors.**
- **Storage tradeoffs**
 - **Local storage vs Distributed storage**
 - **Strong consistency vs Eventual consistency**
- **Storage options are: File storage, Local disk storage, Network attached storage (NAS), Cloud file systems (S3/Blob), Block storage, RAID, Storage area network (SAN), Object storage, HDFS, Streaming storage.**



Batch processing vs Stream processing

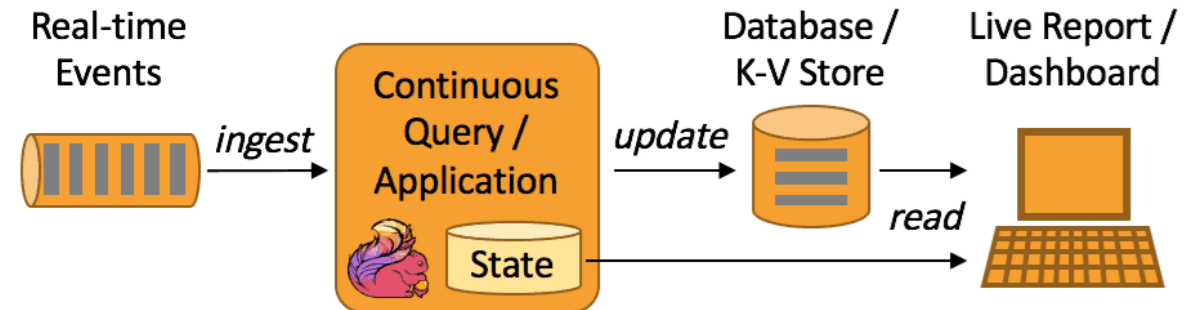
- Processing finite set of data (data at rest).
- Incremental data load is managed by programmer.
- Cluster planned as per data size. High throughput.
- Job run once per batch.

Batch Processing

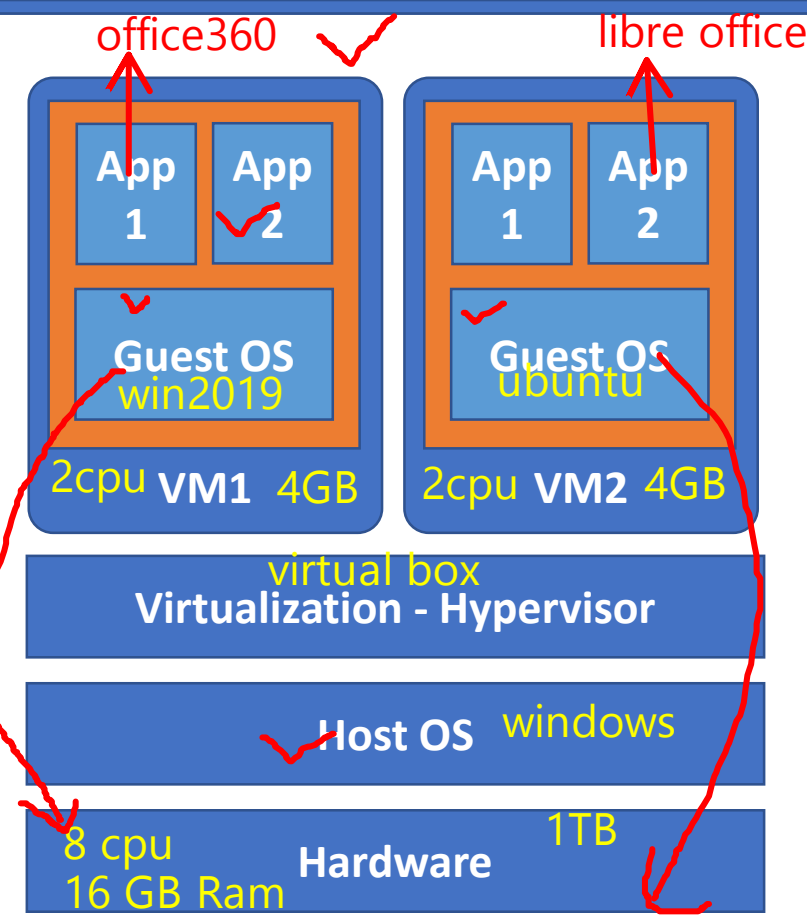


- Processing live stream of data (data in motion).
- Data processing is managed by the framework.
- Less throughput.
- Job is running forever.

Stream Processing




Virtualization vs Containerization



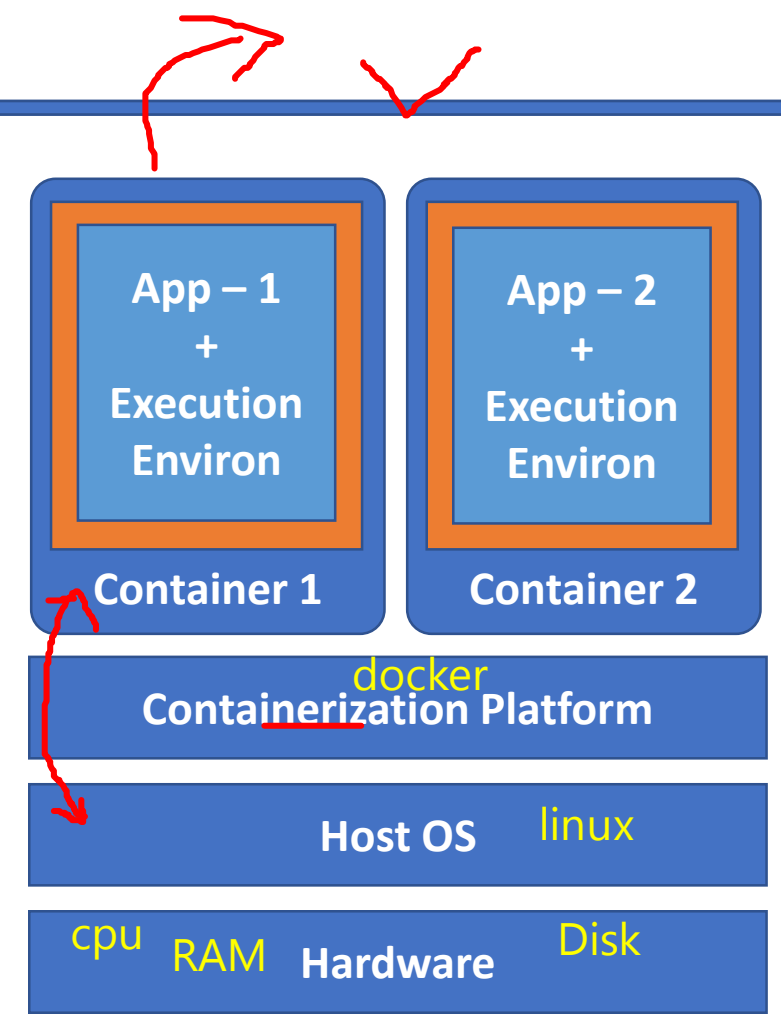
Type-II Virtualization

VMWare, VirtualBox, KVM, ...



Type-I Virtualization

VMWare ESX, XEN, Hyper-V, ...




Containerization

Docker, Podman, rkt, ...



Orchestration

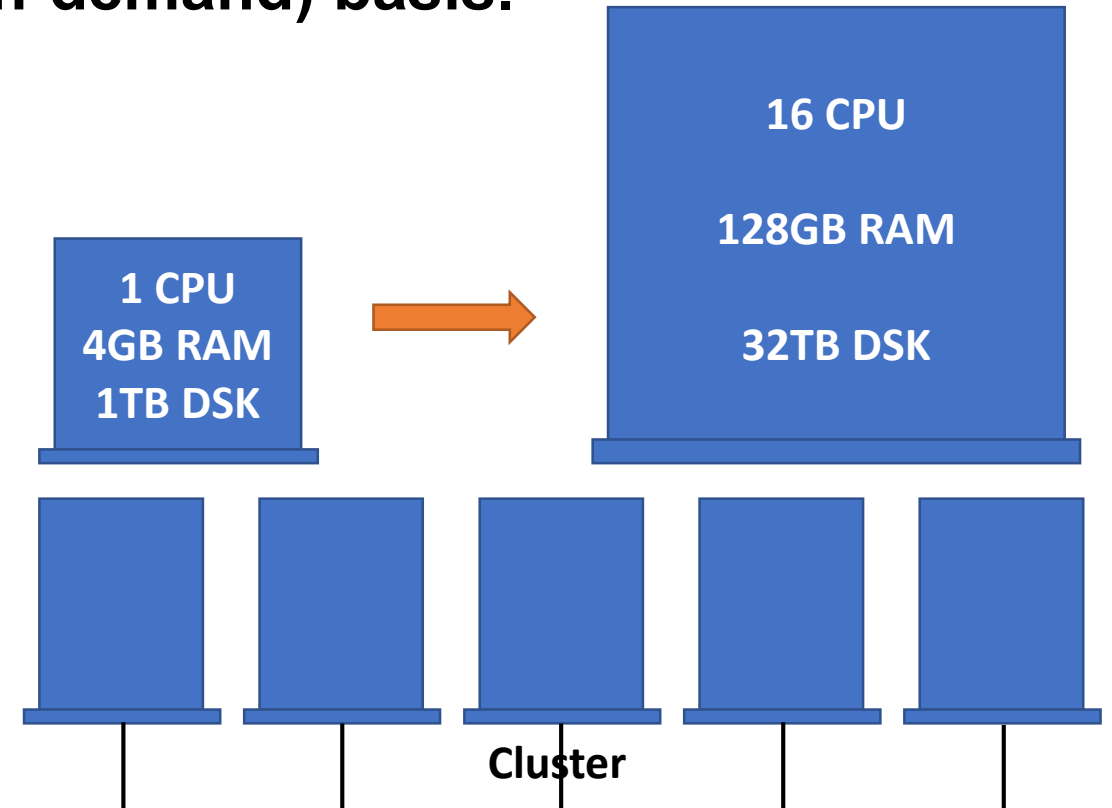
- **Container Orchestration** auto increase or decrease containers to handle change in workloads/demands. It also handles container failure (re-start).
- **Ex: Docker swarm, Kubernetes, ...**



Scalability and Elasticity

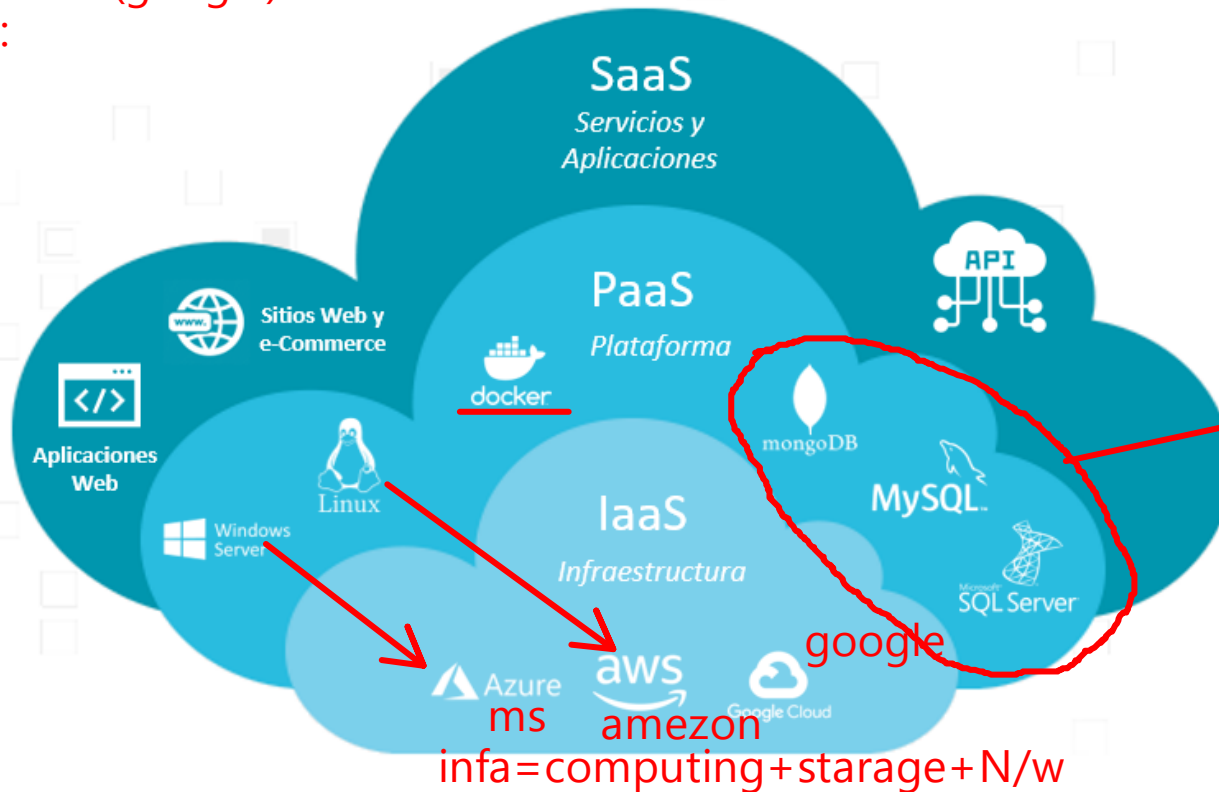
- Scalability is “ability of system / application to perform well under an increased or expanding workload”.
- The resource usage is increased or decreased as per workload.
- Vertical scalability / Up scaling:
 - Increasing single system (hardware) resources in order to handle higher loads.
 - Need to handle SPOF (single point of failure) by adding backup system.
- Horizontal scalability / Out scaling:
 - Adding new systems/nodes into the cluster in order to handle higher loads.
 - More economical solution with higher complexity.

- Elastic: Cloud systems are designed to increase/decrease load as per workload.
- Cloud payments are usually pay-per-use (on-demand) basis.



Cloud Service Models

cloud provider
1: AWS (Amezon)
2: Azure (MC)
3: GCP(google)
4:



- **IaaS: Infrastructure as-a Service**
 - AWS EC2, S3, VPC
- **PaaS: Platform as-a Service**
 - Beanstalk, SageMaker,
- **SaaS: Software as-a Service**
 - Gmail, Drive, Facebook, LinkedIn, Netflix
- **DaaS: Database as-a Service**
 - RDBMS, Aurora, Atlas, DynamoDb AWS
- **FaaS: Function as-a Service**
 - Lambda, Google functions



Big Data & Analytics Spectrum


• Data storage

- RDBMS & NoSQL databases
- Data warehouse
- S3, DFS, ...



• Data Analysis & visualizations

- Data Visualizations
- Business reports





• Artificial Intelligence, Data Science & Data mining

- Mathematics, Statistics & Computer algorithms
- Machine learning & Deep learning
- R Programming, Python



• (Big) Data Engineering

- Hadoop, Hive, Spark, Kafka, BigTable, ...
- Java, Scala, Python.

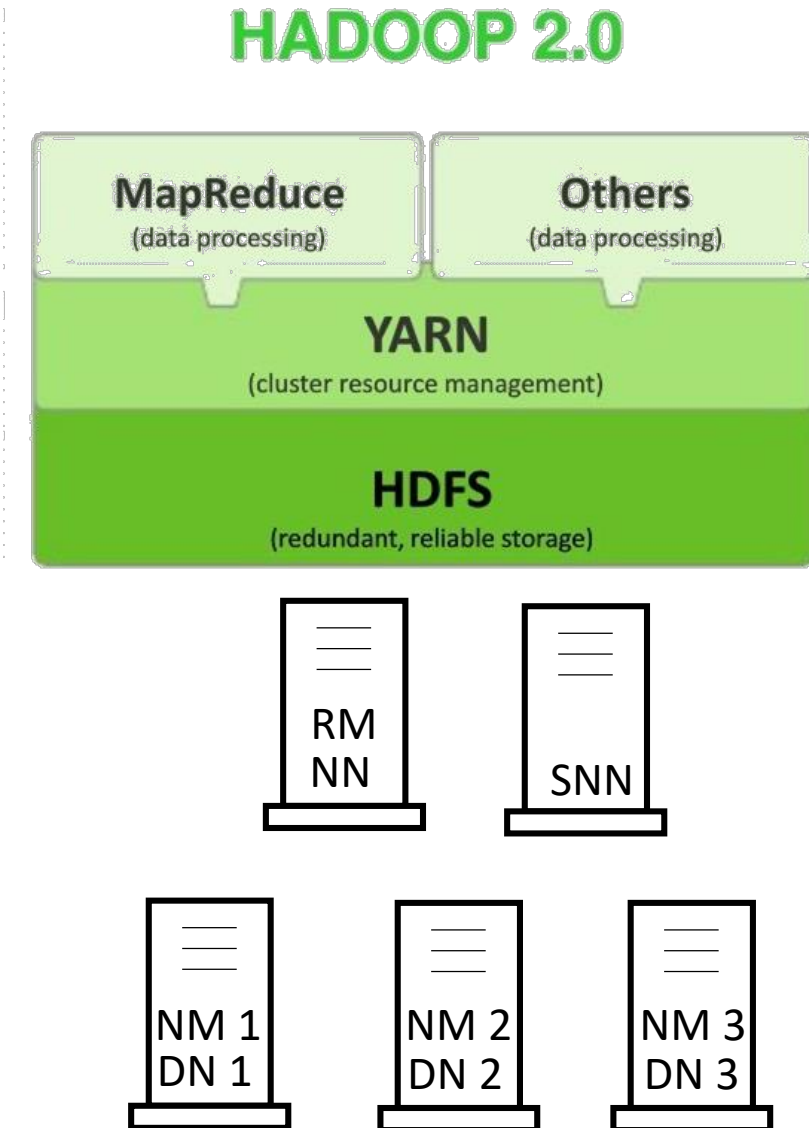



• Infrastructure

- Linux, Cloud Computing




- Hadoop is developed by Doug cutting.
 - Web crawler – Nutch
 - Distributed computing and storage needed to process huge data produced by the crawler.
 - Joined Yahoo. Developed and open sourced under Apache license. @2006
- Hadoop
 - Distributed storage: HDFS
 - Distributed computing Map-reduce
 - Cluster manager: YARN java yet another resource negotiator
- Hadoop is like a Kernel/Platform on which many different applications are built (eco-systems).



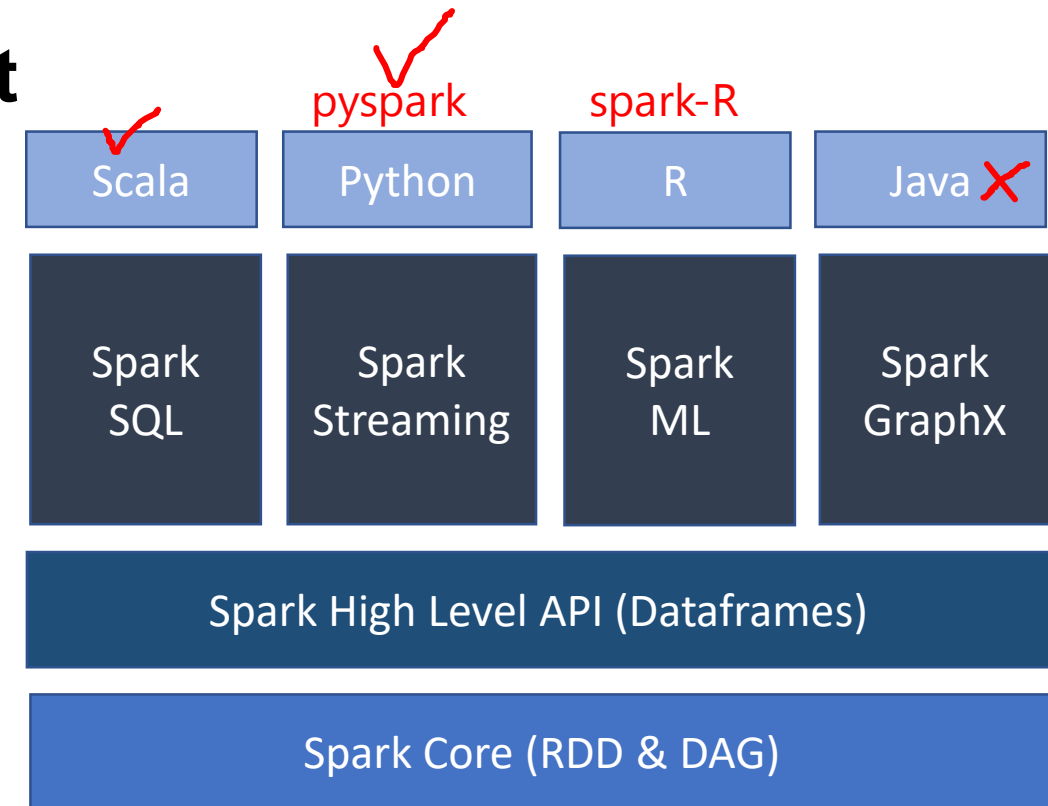
Apache Hive

- Developed by Facebook (2007)
- Client software that convert Hive QL queries to MapReduce.
- Hive QL is similar to SQL with many extended features.
- Hive manages structured data.
- Hive is data warehouse (OLAP) built for Hadoop.
 - Data storage = HDFS
 - Metadata = RDBMS
 - Data processing = Map-reduce or Spark or Tez.



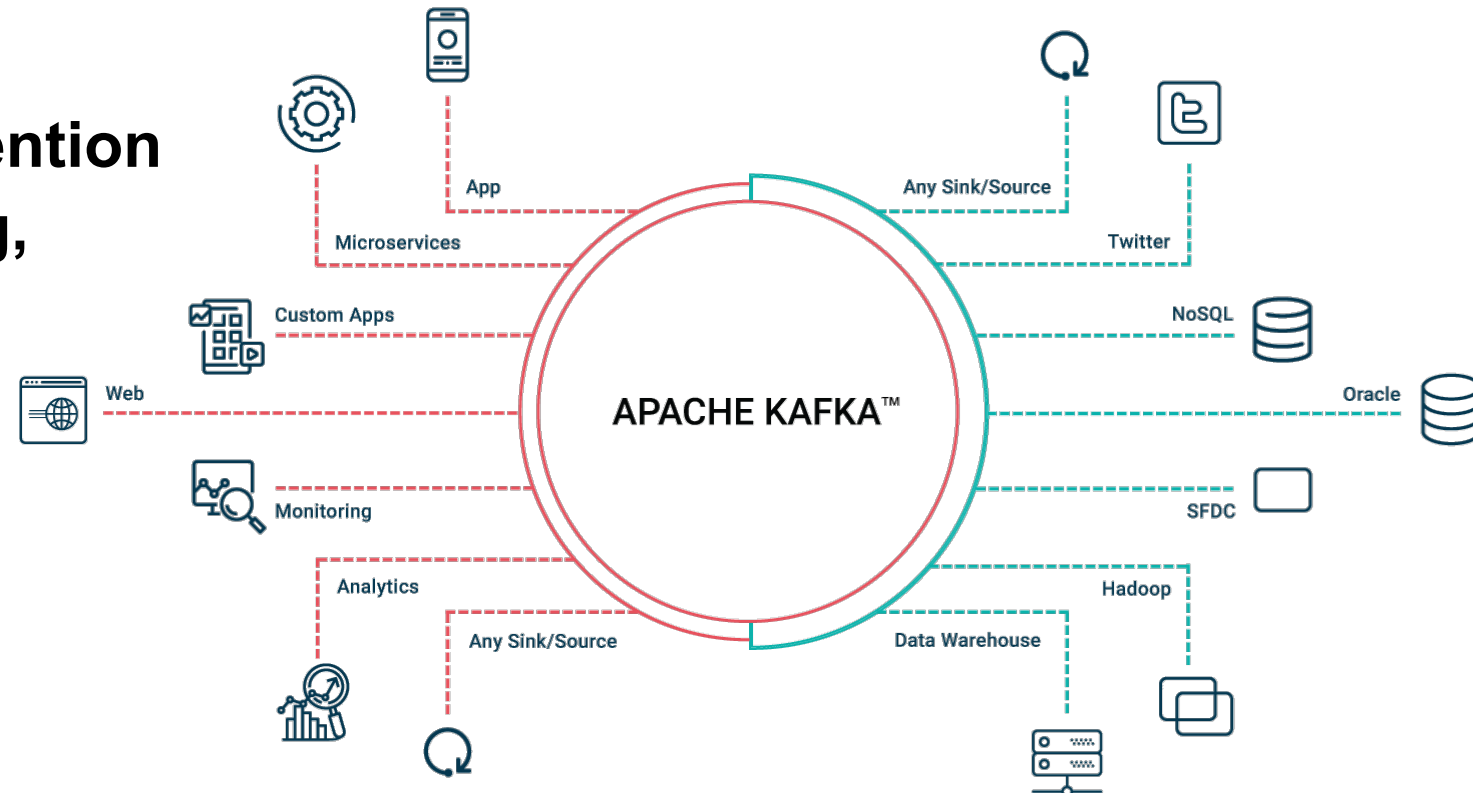
Apache Spark

- Spark is Distributed computing framework, that can process huge amount of data.
- Spark can be used as eco-system of Hadoop or can be used as independent distributed computing framework.
- Developed by UCB AMPLabs division.
- Further developed/maintained by DataBricks.
- Popular Spark vendors
 - DataBricks, AWS EMR, Cloudera, MapR
- Spark Toolkit



Apache Kafka

- Kafka is a distributed messaging system.
- Developed at LinkedIn and open sourced in 2011.
- Used by LinkedIn, Twitter, Uber, airbnb, ...
- **Advantages**
 - Scalable, Durable, Finite retention
 - Low latency, Strong ordering,
 - Exact once delivery
- **Applications**
 - Stream processing
 - Notifications.



Big Data domains & opportunities

- **Domains:** Health-care, Retails, Trading/Share market, Finance, Security, Fraud, Search engines, Log Analysis, Telecom, Traffic Control, Manufacturing and lot more.
- **Big Data is all about :-** Think, Collect, Manage, Analyze, Summarize, Visualize, Discover Knowledge and Take Decisions.

- **Job profiles:**

- Business Analyst/Intelligence
- Database engineer / DWH
- Big Data engineer
- Data operations
- Big Data Architect



Que : Which of the following is used for live data Processing?

- A. Batch**
- B. Stream**
- C. Unit**
- D. Query**



Que : Which of the following is used for live data Processing?

A. Batch

B. Stream

C. Unit

D. Query



Que : In Cloud computing _____ type of Virtualization is used.

- A. Type I**
- B. Type II**
- C. Type A**
- D. Type B**



Que : In Cloud computing _____ type of Virtualization is used.

- A. Type I**
- B. Type II**
- C. Type A**
- D. Type B**



Que : PaaS stands for _____

- A. Python as-a Service**
- B. Platform as-a Service**
- C. Program as-a Service**
- D. Process as-a Service**



Que : PaaS stands for _____

- A. Python as-a Service**
- B. Platform as-a Service**
- C. Program as-a Service**
- D. Process as-a Service**



Que : Which of the programming language is not used for big data processing?

- A. Python
- B. Java
- C. Scala
- D. mongo



Que : Which of the programming language is not used for big data processing?

A. Python

B. Java

C. Scala

D. mongo



Que : Which of the following are incorrect Big Data Technologies?

- A. Hadoop
- B. Spark
- C. Hive
- D. Splunk



Que : Which of the following are incorrect Big Data Technologies?

A. Hadoop

B. Spark

C. Hive

D. Splunk





Thank you!

Pradnyaa S Dindorkar <pradnya@sunbeaminfo.com>

